

# Week 12 Milestone Report

Graph Analytics and Centrality Report  
MovieLens Discovery System — Semester Project

Grupo 06 — Big Data  
BD-Grupo-06/movielens-discovery

July 4, 2026

## Contents

|           |                                                                           |           |
|-----------|---------------------------------------------------------------------------|-----------|
| <b>1</b>  | <b>Executive Summary</b>                                                  | <b>2</b>  |
| <b>2</b>  | <b>Objective</b>                                                          | <b>2</b>  |
| <b>3</b>  | <b>Inputs and Reproducible Artifacts</b>                                  | <b>3</b>  |
| 3.1       | Input artifacts from previous weeks . . . . .                             | 3         |
| 3.2       | Week 12 output artifacts . . . . .                                        | 3         |
| <b>4</b>  | <b>Graph Definition</b>                                                   | <b>3</b>  |
| <b>5</b>  | <b>Graph Construction and Sensitivity Sweep</b>                           | <b>4</b>  |
| <b>6</b>  | <b>Chosen Configuration and Validity Checks</b>                           | <b>5</b>  |
| <b>7</b>  | <b>Centrality Measures</b>                                                | <b>6</b>  |
| <b>8</b>  | <b>Comparison: Graph Ranking vs. Popularity vs. Model-Based Ranking</b>   | <b>7</b>  |
| <b>9</b>  | <b>Interpretation Note: What Graph Structure Means, and Does Not Mean</b> | <b>8</b>  |
| <b>10</b> | <b>Subgraph Sampling for the 3D Web Visualization</b>                     | <b>8</b>  |
| <b>11</b> | <b>What Worked and What Did Not</b>                                       | <b>9</b>  |
| <b>12</b> | <b>Ethics and Access Note</b>                                             | <b>9</b>  |
| <b>13</b> | <b>Reproducibility</b>                                                    | <b>10</b> |
| <b>14</b> | <b>Conclusion</b>                                                         | <b>10</b> |

# 1 Executive Summary

This report covers Week 12: formalizing the item-item graph induced by the MovieLens catalog and using it for structural analysis. The graph reuses the Week 10 SVD item-factor space (the same signal behind the `svd_collaborative` recommender), so this is a genuine structural view of an existing signal, not an unrelated ad hoc graph.

Table 1: Graph overview and headline validity/comparison numbers

| Metric                                                    | Value                                           |
|-----------------------------------------------------------|-------------------------------------------------|
| Nodes                                                     | 13,176 (movies with $\geq 50$ ratings, i.e. a W |
| Edges                                                     | 155,331 (undirected, weighted by cosine s       |
| Sparsity                                                  | 0.179% of all possible pairs                    |
| Isolated nodes                                            | 29 (0.22%)                                      |
| Connected components                                      | 33                                              |
| Giant component coverage                                  | 13,139 nodes ( <b>99.72%</b> )                  |
| Spearman $\rho$ : popularity vs. graph PageRank           | <b>0.313</b>                                    |
| Spearman $\rho$ : model rec. in-degree vs. graph PageRank | <b>0.578</b>                                    |

The graph is highly connected (one dominant component covering 99.7% of nodes) at the chosen configuration ( $k = 15$  neighbors, minimum cosine similarity 0.5). PageRank correlates only weakly with raw popularity ( $\rho = 0.313$ ), confirming that graph centrality captures something structurally different from “most-rated” — a small number of niche, low-popularity titles are highly central because their rating pattern places them in a densely interconnected similarity neighborhood.

## 2 Objective

The goal of Week 12 was to formalize the domain as a graph and use it for structural analysis, connecting back to the Week 10 recommendation layer rather than building an unrelated graph:

*What does the similarity structure among movies look like as a graph, and does structural centrality tell us something different from popularity or model-based ranking?*

Specific questions addressed:

- What do nodes and edges mean, and why are they defined this way?
- Is the chosen graph configuration ( $k$ , similarity threshold) robust, or a hand-picked lucky run?
- Does PageRank on the similarity graph rank movies differently than popularity or the Week 10 recommender’s in-degree?
- What does high centrality mean in this domain, and — just as important — what does it *not* mean?

### 3 Inputs and Reproducible Artifacts

#### 3.1 Input artifacts from previous weeks

Table 2: Inputs consumed from prior weeks

| Artifact                          | Source  | Purpose                                           |
|-----------------------------------|---------|---------------------------------------------------|
| week10_svd_item_factors.parquet   | Week 10 | 50-dim item factors: node features / edge weights |
| week10_popularity_global.csv      | Week 10 | Popularity ranking (comparison baseline)          |
| week10_svd_recs_top20.parquet     | Week 10 | Recommendation in-degree (model-based baseline)   |
| week07_kmeans_assignments.csv     | Week 7  | Cluster labels for the web viewer                 |
| movies_catalog.parquet, links.csv | Week 3  | Titles, genres, IMDb/TMDb IDs                     |

#### 3.2 Week 12 output artifacts

Table 3: Output artifacts produced in Week 12

| File                                    | Description                                                                        |
|-----------------------------------------|------------------------------------------------------------------------------------|
| week12_graph_meta.json                  | Graph definition, chosen configuration, validity checks                            |
| week12_graph_metrics.parquet            | Per-node degree, weighted degree, PageRank, eigenvector                            |
| week12_sensitivity_sweep.csv/.html/.png | Sweep over $k \in \{5, 10, 15, 20\} \times \text{threshold} \in \{0.3, 0.4, 0.5\}$ |
| week12_popularity_vs_pagerank.html/.png | Popularity rank vs. graph PageRank rank scatter plot                               |
| movie_graph_viz.json                    | Sampled 400-node / 3,334-edge subgraph for the viz                                 |

## 4 Graph Definition

**Node:** a movie with at least 50 ratings in the 25M-rating matrix, i.e. a movie that has a Week 10 SVD item factor (`min_movie_ratings=50` per `week10_svd_meta.json`). Restricting nodes to this set — rather than all  $\approx 62,423$  catalog rows — is a deliberate grain choice: movies below the rating floor have no reliable collaborative signal, so a “similarity” edge to them would be noise, not structure. This yields **13,176 candidate nodes**.

**Edge:** undirected, weighted. Weight = cosine similarity between two movies’ 50-dimensional SVD item-factor vectors — the same vectors and metric used by the Week 10 `svd_collaborative` recommender. An edge exists between  $A$  and  $B$  if  $B$  is among  $A$ ’s top- $k$  nearest neighbors by cosine similarity **or**  $A$  is among  $B$ ’s top- $k$  neighbors (symmetrized), and the similarity clears a minimum threshold.

**Why undirected.** Cosine similarity is symmetric by construction ( $\text{sim}(A, B) = \text{sim}(B, A)$ ), so a directed graph would only be justified by an asymmetric relation. That asymmetric relation (“ $B$  appears in  $A$ ’s top-20 recs but not vice versa”) is kept as a *separate* ranking signal (`model_rec_indegree`) for the comparison in Section 8, rather than being conflated with the edge definition.

**Why weighted and thresholded, not fully dense.** A fully dense weighted graph on 13,176 nodes ( $\approx 86.8\text{M}$  possible undirected pairs) would be neither interpretable nor visualizable, and most pairs have near-zero similarity. Keeping only the top- $k$  neighbors per node above a minimum similarity controls sparsity and keeps only edges that plausibly reflect a real co-rating relationship. Both  $k$  and the threshold are swept in Section 5 rather than chosen by hand.

## 5 Graph Construction and Sensitivity Sweep

Edges are built with `sklearn.neighbors.NearestNeighbors` (`metric='cosine'`, brute force), which computes each node's  $k$  nearest neighbors without materializing the full  $13,176 \times 13,176$  similarity matrix, then symmetrizes and drops pairs below the minimum similarity.

$k$  and the minimum-similarity threshold were swept jointly to check whether the graph's structure is sensitive to the choice, rather than reporting one hand-picked, unexamined configuration (Table 4).

Table 4: Sensitivity sweep: edges, sparsity, isolated nodes, and connectivity by  $k$  and threshold. Chosen configuration in bold.

| $k$       | min. sim.  | Edges          | Sparsity      | Isolated  | Components | Giant frac.   |
|-----------|------------|----------------|---------------|-----------|------------|---------------|
| 5         | 0.3        | 54,494         | 0.063%        | 1         | 2          | 99.99%        |
| 5         | 0.5        | 54,036         | 0.062%        | 29        | 33         | 99.72%        |
| 5         | 0.7        | 41,172         | 0.047%        | 1,903     | 1,957      | 84.56%        |
| 10        | 0.3        | 107,115        | 0.123%        | 1         | 2          | 99.99%        |
| 10        | 0.5        | 105,447        | 0.121%        | 29        | 33         | 99.72%        |
| 10        | 0.7        | 74,813         | 0.086%        | 1,903     | 1,957      | 84.56%        |
| <b>15</b> | <b>0.5</b> | <b>155,331</b> | <b>0.179%</b> | <b>29</b> | <b>33</b>  | <b>99.72%</b> |
| 15        | 0.3        | 158,911        | 0.183%        | 1         | 2          | 99.99%        |
| 15        | 0.7        | 104,760        | 0.121%        | 1,903     | 1,957      | 84.56%        |
| 20        | 0.3        | 210,185        | 0.242%        | 1         | 2          | 99.99%        |
| 20        | 0.5        | 203,980        | 0.235%        | 29        | 33         | 99.72%        |
| 20        | 0.7        | 132,124        | 0.152%        | 1,903     | 1,957      | 84.56%        |


Graph sensitivity to  $k$  and minimum-similarity threshold

Figure 1: Isolated-node count and giant-component fraction as functions of  $k$ , grouped by minimum-similarity threshold.

**Key observation.** For a fixed threshold, the isolated-node count is *constant* across all values of  $k$  (1 at threshold= 0.3, 29 at threshold= 0.5, 1,903 at threshold= 0.7). This is because isolation is driven by whether a node's single closest neighbor clears the threshold at all — if the top-1 neighbor's similarity already falls below the threshold, looking at more neighbors ( $k$ ) cannot produce a qualifying edge. The threshold, not  $k$ , controls connectivity;  $k$  mainly controls how many *additional* edges accumulate once a node is already connected.

## 6 Chosen Configuration and Validity Checks

Threshold= 0.7 strands 1,903 nodes (14.4%) even at  $k = 20$  — too strict, since most real co-rating similarity signal in this 50-dimensional space falls below 0.7. Threshold= 0.3 is already near-fully-connected at  $k = 10$  (99.99% giant component, only 1 isolated node) — too permissive, approaching a near-complete graph and defeating the purpose of a sparse structural graph.

**Chosen configuration:**  $k = 15$ , **min\_similarity**= 0.5. This keeps the giant component covering the large majority of nodes (99.72%) while still discarding most low-similarity pairs (sparsity 0.179%), and matches how the Week 10 **content\_cosine** / **svd\_collaborative** recommenders already treat similarity — same latent space, same metric.

Table 5: Validity checks at the chosen configuration

| Check                    | Result                                                          |
|--------------------------|-----------------------------------------------------------------|
| Edge sparsity            | 0.179% of all possible undirected pairs — sparse by design      |
| Isolated nodes           | 29 / 13,176 (0.22%)                                             |
| Connected components     | 33 (one giant component + 32 small fragments)                   |
| Giant component coverage | 13,139 / 13,176 nodes (99.72%)                                  |
| Sensitivity              | Section 5 — connectivity is threshold-driven, not one lucky $k$ |

## 7 Centrality Measures

Degree, weighted degree, and PageRank are computed on the full graph. Eigenvector centrality is undefined on disconnected graphs (`networkx` raises `AmbiguousSolution`), so it is computed on the giant connected component only, with 0 assigned to the 37 nodes outside it — those nodes have no path into the component’s dominant eigenvector by construction, so 0 is the correct value, not a missing one. Exact betweenness centrality is  $O(V \cdot E)$  and infeasible at this scale (13,176 nodes, 155,331 edges); an approximate betweenness is computed instead using a random sample of 500 source nodes (`networkx`’s native `k` parameter), with edge distance defined as  $1 - \text{similarity}$ .

Table 6: Top 10 movies by PageRank, with degree, weighted degree, and popularity/model rank (out of 13,176)

| Movie                       | PageRank | Degree | Wtd. degree | Pop. rank | Model-rec |
|-----------------------------|----------|--------|-------------|-----------|-----------|
| L’Atalante (1934)           | 0.000265 | 107    | 100.32      | 6,351     |           |
| The Lazarus Effect (2015)   | 0.000261 | 96     | 85.01       | 7,985     |           |
| Viridiana (1961)            | 0.000247 | 101    | 95.56       | 5,592     |           |
| Day of Wrath (1943)         | 0.000237 | 94     | 88.33       | 8,052     |           |
| Perfect Stranger (2007)     | 0.000230 | 82     | 68.94       | 6,044     |           |
| Untraceable (2008)          | 0.000228 | 81     | 66.64       | 5,594     |           |
| The Reaping (2007)          | 0.000226 | 82     | 70.13       | 5,950     |           |
| Pickpocket (1959)           | 0.000223 | 90     | 84.13       | 6,120     |           |
| I Know Who Killed Me (2007) | 0.000223 | 79     | 70.20       | 7,632     |           |
| Devil’s Knot (2013)         | 0.000217 | 77     | 61.85       | 10,766    |           |

**Key finding.** None of the top-10 PageRank movies rank in the popularity top-1,000 — the highest is L’Atalante at popularity rank 6,351 out of 13,176 (bottom half of the catalog). High PageRank here means a movie sits in a densely-interconnected neighborhood of the similarity structure (reachable through many strong-similarity chains), which is a structurally different property from being widely rated.

## 8 Comparison: Graph Ranking vs. Popularity vs. Model-Based Ranking

Three independent rankings over the same 13,176-node set:

- **Popularity ranking:** rank by `rating_count` (Week 10 `week10_popularity_global.csv`) — “how many people rated this.”
- **Model-based ranking:** rank by *recommendation in-degree* — how many times a movie appears inside another movie’s Week 10 SVD top-20 recommendation list. This is a directed, model-derived signal distinct from the graph’s undirected similarity edges.
- **Graph ranking:** rank by PageRank on the Week 12 similarity graph.

Table 7: Spearman rank correlations between the three ranking signals

| Comparison                              | Spearman $\rho$ |
|-----------------------------------------|-----------------|
| Popularity vs. graph PageRank           | <b>0.313</b>    |
| Model rec. in-degree vs. graph PageRank | <b>0.578</b>    |
| Popularity vs. model rec. in-degree     | 0.785           |

Popularity rank vs. graph PageRank rank (Spearman  $\rho=0.313$ )




Figure 2: Popularity rank vs. graph PageRank rank (Spearman  $\rho = 0.313$ ). Points scatter widely around the diagonal, showing the two rankings diverge substantially.

Popularity and the model’s recommendation in-degree are strongly correlated ( $\rho = 0.785$ ) — both are, in different ways, proxies for how mainstream a movie is. Graph PageRank correlates much more weakly with popularity ( $\rho = 0.313$ ) and moderately with model in-degree ( $\rho = 0.578$ ): the similarity graph surfaces a partially different

signal — structural embeddedness in the SVD similarity space — that is only loosely tied to how many people rated a movie.

## 9 Interpretation Note: What Graph Structure Means, and Does Not Mean

**What it means:** an edge between two movies indicates that, in the space learned by the Week 10 SVD factorization of the user-rating matrix, the two movies have highly correlated rating patterns across the user base — people who rate one in a particular direction tend to rate the other similarly. High PageRank means a movie sits in a densely-interconnected neighborhood of this similarity structure: not just similar to a few movies, but reachable through many short chains of strong similarity.

**What it does not mean:**

- It is **not** a claim about shared plot, cast, or genre — the graph is built purely from rating co-occurrence, not content features. Two movies can be structurally central neighbors while sharing no genre tag.
- High PageRank is **not** the same as “popular” —  $\rho = 0.313$  is positive but far from 1: the graph surfaces movies that are structurally embedded in the similarity space without necessarily being the most-rated movies overall, and vice versa.
- The graph is **not directed** and does not encode “recommend  $A$  because of  $B$ ” — that asymmetric signal is captured separately by the Week 10 recommendation lists (`model_rec_indegree`) used for comparison above.
- The similarity signal inherits whatever rating biases exist in MovieLens (predominantly English-language, enthusiast-skewed userbase); it should not be read as a universal notion of movie similarity.
- Isolated nodes and small disconnected components are not “unrelated” movies in general — they are movies whose rating pattern did not clear the chosen similarity threshold against any node’s top-15 neighbors, which can also reflect a small or unusual rater population for that title.

## 10 Subgraph Sampling for the 3D Web Visualization

Drawing all 13,176 nodes in a browser-based 3D scene is neither readable nor performant, so a **connected** subgraph of 400 nodes was sampled for the interactive viewer ([web/](#)) using hub-expansion sampling: start from the 40 highest-PageRank nodes in the giant component (structural hubs), then greedily grow the sample by always adding the unselected node with the strongest edge weight to an already-selected node. Because every added node is attached by construction, the sampled subgraph has **no isolated nodes**, unlike a uniform-random node sample with an induced-edge cut, which typically strands many isolated points.



Table 8: Full graph vs. sampled subgraph used by the web viewer

|                | Full graph | Sampled subgraph    |
|----------------|------------|---------------------|
| Nodes          | 13,176     | 400                 |
| Edges          | 155,331    | 3,334               |
| Isolated nodes | 29         | 0 (by construction) |

The sampled subgraph is exported to `artifacts/week12/movie_graph_viz.json` and copied to `web/public/movie_graph.json`, where node metadata (title, genres, IMDb/TMDb links, PageRank, degree, cluster ID) drives the interactive viewer. Poster URLs are filled in separately by `scripts/fetch_posters.py` via the TMDb API, for exactly this sampled set of movies.

## 11 What Worked and What Did Not

### What worked.

- Reusing the Week 10 SVD item-factor space as the graph’s edge weight keeps the graph analysis connected to an existing, already-validated signal instead of introducing an unrelated ad hoc similarity metric.
- The sensitivity sweep cleanly separates the role of  $k$  (marginal edge accumulation) from the role of the similarity threshold (connectivity), which made the configuration choice defensible rather than arbitrary.
- Hub-expansion sampling produced a fully connected, hub-anchored 400-node subgraph suitable for a real-time 3D viewer, with zero isolated nodes.

### What did not work.

- **Eigenvector centrality** is undefined outside the giant component; 37 nodes necessarily receive a 0 that reflects “no path into the dominant eigenvector,” not “zero centrality” in an absolute sense — this distinction has to be stated explicitly or it reads as a missing value.
- **Betweenness centrality** could not be computed exactly at this scale (13,176 nodes, 155,331 edges); the 500-source approximation trades precision for tractability and should not be read as an exact ranking.
- **Weak correlation with popularity** ( $\rho = 0.313$ ) means graph PageRank alone is a poor stand-in for “movies most people already know” — it answers a different question (structural embeddedness), not “what should the homepage show a new user.”

## 12 Ethics and Access Note

- **Data source:** GroupLens MovieLens 25M public release (research license).
- **Access:** Educational and research use; raw data not redistributed.
- **Personal data risk:** all user identifiers are anonymized integer IDs.

- **Mitigation:** the graph is built entirely at the movie level from aggregate SVD item factors; no individual user rating history or user identifier is exposed by any node, edge, or centrality score in this analysis.

## 13 Reproducibility

Run the notebook after the Week 10 artifacts exist:

1. `notebooks/week12/week12_graph_analytics.ipynb`

Then, to refresh the web viewer's data:

2. `python scripts/fetch_posters.py` — populates `posterUrl` for the  $\approx 400$  sampled movies via the TMDb API.
3. Copy `artifacts/week12/movie_graph_viz.json` to `web/public/movie_graph.json`.

Table 9: Key parameters for reproducibility

| Parameter                         | Value                                                        |
|-----------------------------------|--------------------------------------------------------------|
| Node definition                   | Movie with $\geq 50$ ratings (has a Week 10 SVD item factor) |
| Candidate nodes                   | 13,176                                                       |
| Similarity metric                 | Cosine, on 50-dim Week 10 SVD item factors                   |
| Chosen $k$ (neighbors per node)   | 15                                                           |
| Chosen minimum similarity         | 0.5                                                          |
| Betweenness sample size           | 500                                                          |
| Subgraph sample size (web viewer) | 400 nodes / 3,334 edges                                      |
| Hub seeds for sampling            | 40 (highest PageRank in giant component)                     |
| Random state                      | 42                                                           |

## 14 Conclusion

Week 12 is complete. A 13,176-node, 155,331-edge item-item similarity graph was built directly from the Week 10 SVD item-factor space, validated with a joint sensitivity sweep over  $k$  and similarity threshold, and analyzed with degree, weighted degree, PageRank, eigenvector, and approximate betweenness centrality.

- The chosen configuration ( $k = 15$ , `min_similarity= 0.5`) yields a giant component covering 99.72% of nodes, and the sweep shows this connectivity is driven by the threshold, not a lucky choice of  $k$ .
- Graph PageRank correlates only weakly with popularity ( $\rho = 0.313$ ) and moderately with the Week 10 model's recommendation in-degree ( $\rho = 0.578$ ): structural centrality in the similarity graph is a genuinely different signal from “most-rated” or “most-recommended,” not a relabeling of either.
- The interpretation note makes explicit what the graph does and does not license: it reflects rating co-occurrence, not content similarity; it is undirected and does not encode asymmetric recommendation; and it inherits MovieLens's own rating-population biases.

- A connected, hub-anchored 400-node subgraph was sampled and exported to drive the interactive 3D MovieLens Discovery viewer ([web/](#)), completing the pipeline from raw ratings through representation, clustering, recommendation, and now graph structure.